

NOVA

IMS

Information
Management
School

BUSINESS INTELLIGENCE I COURSE'S PROJECT

REPORT ON DELIVERY 1

Professor Pedro Batista

Group 049

Raimundo Mujica Costa M20221342

Simão Pedro Acciaioli Gouveia Dos Santos M20220561

Tomás Caldeira Cardoso Soares Esteves M20220377

Waleed Bin Tanveer Sabir M20220073



Table of Contents

Introduction	3
The Company	4
Introduction to Killer Glute Bikes (KGB) company	4
Business Model	4
Problem Scenario and Business needs	5
Data Sources	6
Killer Glute Bikes (OLTP Database)	6
Data Entities and their Attributes	6
Data Entities' Relations	10
Developing the Data Warehouse	11
Getting the right based approach on design the Data Warehouse to answer Business Needs	12
Slowly Changing Dimensions	14
Assembled Data Warehouse: Dimension Tables and Schema	15
Dimension Date	15
Dimension Product	16
Dimension Territory	16
Dimension Store	17
Dimension Costumer	17
Dimension Sales	18
Data warehouse schema	19
Hierarchies	20
Date Hierarchy	20
Product Hierarchy	21
Customer Hierarchy	21
Store Hierarchy	22
Territory Hierarchy	22
Conclusion on Phase 1	Error! Bookmark not defined.

Table of Illustrations

FIGURE I SUPPLY CHAIN MODEL	4
FIGURE II DATABASE TABLES' RELATION	10
FIGURE III BUSINESS QUESTIONS REQUIRED MEASURES	12
FIGURE IV ENTITIES AND THEIR DIMENSIONS	13





FIGURE V SCD TYPES	14
FIGURE VI DIMENSION DATE TABLE	15
FIGURE VII DIMENSION PRODUCT TABLE	16
FIGURE VIII DIMENSION TERRITORY TABLE	16
FIGURE IX DIMENSION STORE TABLE	17
FIGURE X DIMENSION COSTUMER TABLE	17
FIGURE XI DIMENSION SALES TABLE	18
FIGURE XII DATA WAREHOUSE SCHEMA	19
FIGURE XIII DATE HIERARCHY	20
FIGURE XIV PRODUCT HIERARCHY	21
FIGURE XV COSTUMER HIERARCHY	21
FIGURE XVI STORE HIERARCHY	22
FIGURE XVII TERRITORY HIERARCHY	22





Introduction

In the advent of the digital age, the data we collect has surpassed the amount we predicted. This change has caused businesses to evolve and adapt to this new era. With the ever-increasing competition between companies, any edge provided is very beneficial.

This advantage can be had in many ways than one. Data has taken its rightful place as one of the most powerful tools at a company's disposal. Gaining an advantage using data can include but is not limited to streamlining logistics to decrease turnover time to the customer to seeing what products generate a greater return in a much more quantifiable way. Managing a sizeable company becomes a lot easier whilst also providing better strategies for the future of the company, employees, and customers.

One of the key challenges amongst many in gaining an advantage using data is to extract valuable information from a sea of unfiltered data. There are tools that are very efficient in collecting data but in most scenarios, this data is not actionable evidence towards a future strategy. There are systems that can be implemented to filter out the necessary information. The next step is storing this and analysing the useful information.

The potential of these steps is being realized by many business owners and there's a natural evolution in using Business Intelligence solutions to gain the maximum advantage against the competition. Storing and analysing data is one of the early steps in creating a system that will aid the business in making better decisions based on information gathered. This is further compounded with identifying relevant markers that provide valuable insight into the market in which the company operates.

This structured approach to managing a business with the help of using and organizing data is referred to as Data Warehousing (DW). All the information is sorted and divided into multiple streams and used by relevant departments.

During Business Intelligence I, we were tasked with organizing and creating a DW solution for Killer Glute Bikes. We intend to streamline their data flow by organizing and sorting their client and company data followed by loading it into an intelligent solution (ETL). This will help in making informed short-term decisions for the company whilst also paving a way for future strategies that will pay dividends in the long run and prove to be the determining factors in the business's long-term growth and success.



The Company

Introduction to Killer Glute Bikes (KGB) company

“At Killer Glute Bikes (KGB) we believe in a healthy lifestyle for you and your family whilst facilitating an unforgettable outdoors experience for all!

We believe that it is in the wild where one finds their true self and helps awaken a satiation for nature and all its gifts.

Since 1905, we have been your international outdoor companion to aid your experience in the true transformational power that nature has to offer. We offer top quality bikes, gear and apparel coupled with our expert advice to help you create inspiring stories and experience special moments alone or with your family and friends. And because we do not have shareholders, your proceeds with KGB will directly benefit the outdoors and support a sustainable business in the fight for life outdoors.

So, whether you are new to the outdoors or a seasoned pro, we hope you will join us on this journey!”

Business Model

Killer Glute Bikes is a worldwide bike and bike apparel seller that operates in three different continents (Europe, Oceania, and North America), five different countries. For that reason, to grant the same quality service globally and therefore maintaining its position as a leading company within its line of business, Killer Glute Bikes structure follows some specific guidelines that ease this mission.

Each product that is available at the KGB stores comes from the same exact supplier. Allying this principle with a meticulous choice in their providers guarantees that Killer Glutes Bikes can deliver the same top tier quality products in all its stores. Even that this model might increase the freight in some products supply chain, with this business plan the company does not have to spend any resources on external warehouses fees.




Figure 1 Supply Chain model

Killer Glute Bikes company organization found that the best way to keep its promise of delivering a specialized advice on selling their products to their clients was to fill all its stores with at least one expert in bikes’ related subjects.



Problem Scenario and Business needs

KGB is currently undergoing a restructuring of its organizational management process after a purchase of most of its shares by a new owner. As such, the new CEO needs to understand the current situation of the company.

Having a Data Warehouse with all the information compiled will allow a more stoic approach to solve company's problems and get a better understanding of the business itself.

For the new board, the primary focus is to optimize the company's spendings in the multiple markets they are committed. They hope that with the practice of more sustainable procedures, the organization will be more improve their logistical systems to decrease turnover time and increase the understanding of sales to make better decisions for future strategies pertaining to products.

To properly understand the company's reality and figuring out its next move, we identified the following questions:

1. What is the type of products most sold by market and date?
2. What is the average transportation cost, per type of transport, per product, per Market?
3. What is the client segment with the biggest sales by market and date?
4. What is our customer profile by product, product category, and product subcategory?
5. How do product sales vary during the year?
6. What is our average revenue per sale per product type?
7. Which region has higher profits? What is the best investment resolution?





Data Sources

Killer Glute Bikes (OLTP Database)

The data of Killer Glute Bikes company describes their sales, shippers' activity and covers some general facts that, although not being directly involved in the business itself, provide crucial details for the company's operations. For assembling this database, we considered only "meaningful data". This means only the data that we initially perceived as being useful on the resolution of the company's problem scenario was extracted from Killer Glute Bikes' huge dataset.

Data Entities and their Attributes

We selected the entities and their attributes which we considered to give us some good insights on subjects we expected to be relevant on a furtherer analysis.

■ **Date table**

A date table is crucial to carry on this data analysis as it provides an extensive cover on time dimensions, essential to a transactional database.

String type attributes:

- full_name_date [nvarchar(100) not null]
- day_name [nvarchar(50) not null]
- month_name [nvarchar(50) not null]

Numeric type attributes:

- 🔑 pk_date [int not null] (yyyymmdd code)
- day_number_of_week [int not null]
- month_number_of_year [int not null]
- day_number_of_year [int not null]
- week_number_of_year [int not null]
- calendar_quarter [int not null]
- calendar_year [int not null]
- day_number_of_month [int not null]

Date and Time type attributes:

- bk_period [date not null]

■ **ProductCategory table**

Data about the category of the product.

String type attributes:

- 🔑 pk_productcategory [nvarchar(20)]
- english_productcategory_name [nvarchar(50)]
- german_productcategory_name [nvarchar(50)]
- french_productcategory_name [nvarchar(50)]





■ **ProductSubCategory table**

Data about the subcategory of the product. Provides a more precise partition on products' type.

String type attributes:

- 🔑 pk_productcategory [nvarchar(50)]
- fk_productcategory [nvarchar(20)] ref *pk_productcategory in Product Category Table*
- english_productsubcategory_name [nvarchar(200)]
- german_productsubcategory_name [nvarchar(200)]
- french_productsubcategory_name [nvarchar(200)]

■ **Suppliers table**

Data regarding the suppliers of the products.

String type attributes:

- 🔑 pk_supplier [nvarchar(20)]
- supplier_name [nvarchar(50)]
- contact_details [nvarchar(100)]

Numeric type attributes:

- transportation_fees [int not null] (if the fees for transportation are charged)

■ **Products table**

Detailed data regarding products' specifications.

String type attributes:

- fk_productcategory [nvarchar(50) not null] ref *pk_productcategory in Product Category Table*
- fk_productsubcategory [nvarchar(50) not null] ref *pk_productsubcategory in Product SubCategory Table*
- model_name [nvarchar(50)]
- details [nvarchar(250)]
- fk_supplier [nvarchar(50)]

Numeric type attributes:

- 🔑 pk_productkey [nvarchar(20)]
- standard_cost [decimal(18,10)]
- colour [nvarchar(50)]
- list_price [decimal(18,10)]
- size [int]
- active [int]

Date and Time type attributes:

- product_start_date [date not null]
- product_end_date [date not null]





■ **Territory table**

General geographical data.

String type attributes:

- 🔑 pk_territorykey [nvarchar(50)]
- region [nvarchar(50)]
- country [nvarchar(50)]
- group_division [nvarchar(50)]

■ **SubTerritory table**

Geographical data that complements Territory table.

String type attributes:

- 🔑 pk_territorysubcategory [nvarchar(50)]
- fk_territorykey [nvarchar(50)] ref *pk_territorykey in Territory Table*
- city_name [nvarchar(50)]
- postal_code nvarchar(50)

Numeric type attributes:

- bike_lanes_km [decimal(18,10) not null]
- protected_bike_lanes_km [decimal(18,10) not null]

■ **Stores table**

Data about the company's stores network.

String type attributes:

- 🔑 pk_storeid [nvarchar(50)]
- fk_territorysubcategory [nvarchar(50)] ref *pk_territorysubcategroy in SubTerritory Table*
- city [nvarchar(50)]

Numeric type attributes:

- area_m2 [decimal(18,10) not null]
- rent_per_month [decimal(18,10)]
- staff_instore [int not null]
- expert_staff [int not null]
- average_staff_salary_month [int not null]
- accessibility_1to5 [int not null] (measures the level of accessibility from category one to category five)
- showcase_quality_level_1to5 [int not null] (measures the level of the showcase from category one to category five)





■ **Customers table**

Company's client data.

String type attributes:

- customer_name [nvarchar(50)]
- marital_status [nvarchar(50)]
- gender [nvarchar(50)]
- occupation [nvarchar(50)]
- address_line [nvarchar(50)]
- phone [nvarchar(50)]

Numeric type attributes:

- 🔑 pk_customerkey [int]
- yearly_income [int not null]
- number_children_at_home [int not null]
- number_cars_owned [int not null]
- fk_date_first_purchaseid [int] ref *pk_date in Date Table*

Date and Time type attributes:

- birth_date [date not null]

■ **FactSales table**

Sales and shipment data.

String type attributes:

- 🔑 fk_territorykey [nvarchar(50)]
- 🔑 fk_territorysubcategory [nvarchar(50)]
- 🔑 fk_storeid [nvarchar(50)]

Numeric type attributes:

- 🔑 fk_date [int not null]
- 🔑 fk_productkey [int not null]
- 🔑 fk_customerkey [int not null]
- sales_order_line_number [int not null]
- order_quantity [int not null]
- unit_price [decimal(15,3) not null]
- product_standard_cost [decimal(10,3) not null]
- sales_amount [decimal(15,5) not null]
- tax_amount_unit [decimal(10,3) not null]
- tax_amount [decimal(15,3) not null]
- freight [decimal(10,5) not null]
- total_cost [decimal(15,5) not null]

Date and Time type attributes:

- order_date [date]



Data Entities' Relations

We created a diagram in SQL Management Studio that can show the dataset connections with great clarity.

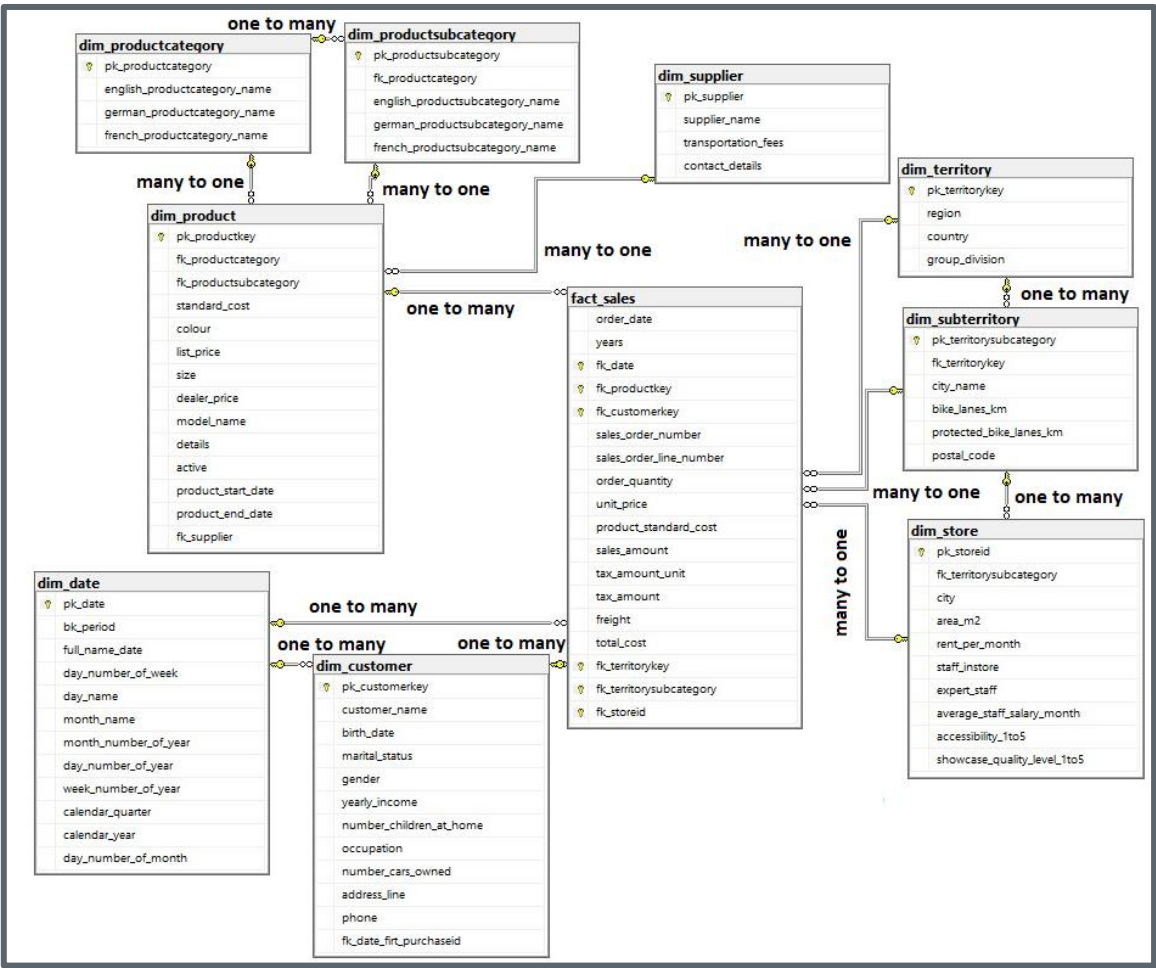


Figure 4 Database Tables' Relations



Developing the Data Warehouse

The main goal of a data warehouse is to store data in an appropriate way so that it is the major support on decision-making process. The principal characteristics that define a data warehouse are the following:

- **Subject-oriented:** as stated before, the main purpose of the data warehouse is to give a support to decision makers by analysing data. It is not designed to assist on daily operations.
- **Integrated:** the assemble of a data warehouse is based upon integrating heterogenous sources.
- **Time-Variant:** the data stored in a data warehouse is not keeping up with the company's live data. Instead, this data provides information from an historical point of view.
- **Non-volatile:** the data warehouse and the data used in the operational environment are always set up individually.

We envisioned to use a Star Schema design for our Data Warehouse as despite using a considerable amount of space, it allows an easier and more concise information split. On the following subchapters we describe our thinking process and appoint the reasons that explain the Killer Glute Bikes' data warehouse design.

Methodology

After assuring every team member had the right mindset and impressions on the importance and requisites for the under development Wonderful Wines of the World Data Warehouse, we developed a working task procedure to ensure we reach our goals.

First, bearing in mind the business case questions we need to approach, we would check if we had the necessary measures to evaluate WWW's situation. Probably, after this check-up we would have to define new attributes that will aid the Data Warehouse's users to reach the desired conclusions. We also went through a second look on all of data we had at our disposal: this would give us an indispensable assistance on reaching the Data Warehouse's design that better fits the company's purpose.

This breakthrough scanning will prove to be essential when we finally "decide on the looks" of the data warehouse. At this point, the decision to keep the warehouse as simple and easily readable to its users and future re-developers was already taken.

After structuring what we thought as an efficient assembly line for the pre-schematized Data Warehouse, we finally began to give it form.





Getting the right based approach on design the Data Warehouse to answer Business Needs

To get the right overview of the company's current business reality and achieving a profile on possible resolutions to its future goals, we pointed out the required measures that must figure in the Data Warehouse.

Questions	Measures
What is the type of products most sold by market and date?	order_quantity
What is the average transportation cost, per type of transport, per product, per Market?	freight
What is the client segment with the biggest sales by market and date?	birth_date
What is our customer profile by product, product category, and product subcategory?	marital_status, gender, yearly_income, number_children_at_home, number_owned, occupation
How do product sales vary during the year?	sales_amount
What is our average revenue per sale per product type?	sales_amount
What stores should be closed? Which ones deserve a higher investment?	sales_amount
Which region has higher profits? What is the best investment resolution?	profit

Figure 11: Business questions required measures

The *profit* measure does not figure in KGB's dataset. As it plays a crucial role in further analysis, we will have to create a new field for evaluating this subject.

By using the attributes *sales_amount* and *total_cost*, both from FactSales table we will be able to get this measure.

$$\text{profit} = \text{sales_amount} - \text{total_cost}$$

We also will be asked on counselling on closing stores. For that reason, we will establish an attribute referring to the store status: active or unactive.



We have sorted out the relevant information broken down into multiple dimensions that would be needed in our DW solution for KGB coupled with the OLTP data source which is going to input a regular stream of new data to the relevant dimensions. These dimensions were carefully chosen based on the attributes that are going to be beneficial in the analyzation phase.

Dimensions	OLTP Entities	Key Data Type
sk_date	Dbo.Date	int
bk_period	Dbo.Date	date
full_name_date	Dbo.Date	nvarchar(50)
day_number_of_week	Dbo.Date	int
day_name	Dbo.Date	nvarchar(50)
month_name	Dbo.Date	nvarchar(50)
month_number_of_year	Dbo.Date	int
day_number_of_year	Dbo.Date	int
week_number_of_year	Dbo.Date	int
calendar_quarter	Dbo.Date	int
calendar_year	Dbo.Date	int
day_number_of_month	Dbo.Date	int
sk_productid	n/a	int
bk_product	Dbo.Product	int
bk_productcategory	Dbo.ProductCategory	nvarchar(50)
bk_productsubcategory	Dbo.ProductSubCategory	nvarchar(50)
productcategory_name	Dbo.ProductCategory	nvarchar(50)
productsubcategory_name	Dbo.ProductSubCategory	nvarchar(50)
standard_cost	Dbo.Product	decimal(18, 10)
colour	Dbo.Product	nvarchar(50)
list_price	Dbo.Product	decimal(18, 10)
size	Dbo.Product	int
dealer_price	Dbo.Product	decimal(18, 10)
model_name	Dbo.Product	nvarchar(50)
details	Dbo.Product	nvarchar(250)
product_start_date	Dbo.Product	date
product_end_date	Dbo.Product	date
bk_supplier	Dbo.Supplier	nvarchar(50)
supplier_name	Dbo.Supplier	nvarchar(50)
transportation_fees	Dbo.Supplier	int
sk_customerid	n/a	int
sk_customer	Dbo.Customer	int
customer_name	Dbo.Customer	nvarchar(50)
birth_date	Dbo.Customer	date
marital_status	Dbo.Customer	nvarchar(50)
gender	Dbo.Customer	nvarchar(50)
yearly_income	Dbo.Customer	int
number_children_at_home	Dbo.Customer	int
occupation	Dbo.Customer	nvarchar(50)
number_cars_owned	Dbo.Customer	int
sk_storeid	n/a	int
bk_store	Dbo.Store	nvarchar(50)
area_m2	Dbo.Store	decimal(18, 10)
rent_per_month	Dbo.Store	decimal(18, 10)
staff_instore	Dbo.Store	int
expert_staff	Dbo.Store	int
average_staff_salary_month	Dbo.Store	int
accessibility_1to5	Dbo.Store	int
showcase_quality_level_1to5	Dbo.Store	int
active	n/a	int
sk_territoryid	n/a	int
bk_territorysubcategory	Dbo.TerritorySubCategory	nvarchar(50)
bk_territorykey	Dbo.Territory	nvarchar(50)
region	Dbo.Territory	nvarchar(50)
group_division	Dbo.Territory	nvarchar(50)
country	Dbo.Territory	nvarchar(50)
city_name	Dbo.TerritorySubCategory	nvarchar(50)
bike_lanes_km	Dbo.TerritorySubCategory	decimal(18, 10)
protected_bike_lanes_km	Dbo.TerritorySubCategory	decimal(18, 10)

Figure 3: Entities and their Dimensions



Slowly Changing Dimensions

Optimizing the Data Warehouse for better responding to changes that may occur in the future is essential. For this reason, we consider implementing two different types of SCD:

Type 1 (overwrite)	Dim_Date
	Dim_Customer
	Dim_Territory
Type 2 (add a new record)	Dim_Product
	Dim_Store

Figure 3. SCD types

For Dim_Date, Dim_Customer and Dim_Territory we decided not to implement a new record for future changes, but to rewrite on the previous values. The reason behind this decision comes from the vague modifications that will have an impact on the business decisions over time.

We believe that Dim_Product and Dim_Store have important information about our business needs, and we expect these dimensions to change over time. Implementing a new record such as a SCD will allow us to have complete information from start to finish. In addition, for future analysis, since we have complete information about our business, we will obtain better results. For Dim_Product we added two columns that represent the start date of a product and the date when the same product changed supplier or when the same product stopped producing. For Dim_Store we add a column that shows which are the stores that are active, representing them with a number, one, and which are the stores that are inactive, representing them with a number, zero.

For Dim_Product, as we mentioned before, we consider that some products in the future will suffer changes. One of those changes would be if the product modifies its specification, like supplier. Another of those changes would be if we decide to stop the production of a specific product. Finally, it is probable that in the future we will decide to change the name of a category or subcategory. Adding a new record will help us to have the complete history of our product over time.

In the case of Dim_Store, we do not know what could happened on the KGB's future. For this reason, we must face a situation in which the company does not have the best results and we must close some stores. Add an additional column implementing the SCD type 2 will help us for a situation like this. So, we can have the complete store's status over time.



Assembled Data Warehouse: Dimension Tables and Schema

We decided to establish six dimensions that we found that would manage to answer all Killer Glute Bikes’ questions and needs. The dimensions were assembled on a certain order (figure vi), that ensures the Data Warehouse’s relations existence. We methodize this step by assuring the dimensions with foreign keys were only coded after the pointed attribute was already established. Briefly, we followed the logical granularity for our Data Warehouse’s design.




Figure vi Dimensions' order creation

Dimension Date

This table contains all the information pertaining to the calendar. This table serves to be a reference point for the billing department to keep track of day-to-day sales and provides an outlook on aggregate sales (monthly, yearly, etc.).

	Column Name	Data Type	Allow Nulls
▼	sk_date	int	<input type="checkbox"/>
	bk_period	date	<input type="checkbox"/>
	full_name_date	nvarchar(50)	<input checked="" type="checkbox"/>
	day_number_of_week	int	<input type="checkbox"/>
	day_name	nvarchar(50)	<input checked="" type="checkbox"/>
	month_name	nvarchar(50)	<input checked="" type="checkbox"/>
	month_number_of_year	int	<input type="checkbox"/>
	day_number_of_year	int	<input type="checkbox"/>
	week_number_of_year	int	<input type="checkbox"/>
	calendar_quarter	int	<input type="checkbox"/>
	calendar_year	int	<input type="checkbox"/>
	day_number_of_month	int	<input type="checkbox"/>
			<input type="checkbox"/>

Figure vi Dimension Date Table

Dimension Product

The product table provides the information about all the products the company sells. Including but not limited to the pricing, description, model name etc.

	Column Name	Data Type	Allow Nulls
▶	sk_productid	int	<input type="checkbox"/>
	bk_productkey	int	<input type="checkbox"/>
	bk_productcategory	nvarchar(50)	<input checked="" type="checkbox"/>
	bk_productsubcategory	nvarchar(50)	<input checked="" type="checkbox"/>
	productcategory_name	nvarchar(50)	<input checked="" type="checkbox"/>
	productsubcategory_name	nvarchar(50)	<input checked="" type="checkbox"/>
	standard_cost	decimal(18, 10)	<input checked="" type="checkbox"/>
	colour	nvarchar(50)	<input checked="" type="checkbox"/>
	list_price	decimal(18, 10)	<input checked="" type="checkbox"/>
	size	int	<input checked="" type="checkbox"/>
	dealer_price	decimal(18, 10)	<input checked="" type="checkbox"/>
	model_name	nvarchar(50)	<input checked="" type="checkbox"/>
	details	nvarchar(250)	<input checked="" type="checkbox"/>
	product_start_date	date	<input type="checkbox"/>
	product_end_date	date	<input checked="" type="checkbox"/>
	bk_supplier	nvarchar(50)	<input checked="" type="checkbox"/>
	supplier_name	nvarchar(50)	<input checked="" type="checkbox"/>
	transportation_fees	int	<input type="checkbox"/>

Figure 10 Dimension Product Table

Dimension Territory

The territory table portrays the regions that the company actively operates in. This provides with the necessary information pertaining to which region has more sales in the comparison to the rest. This helps with managing inventory based on the sales region. The dimensions in this table like bike lanes and protected bike lanes provide us with precursors to expanding in newer regions based on these indicators.

	Column Name	Data Type	Allow Nulls
▶	sk_territoryid	int	<input type="checkbox"/>
	bk_territorysubcategory	nvarchar(50)	<input checked="" type="checkbox"/>
	bk_territorykey	nvarchar(50)	<input checked="" type="checkbox"/>
	region	nvarchar(50)	<input checked="" type="checkbox"/>
	group_division	nvarchar(50)	<input checked="" type="checkbox"/>
	country	nvarchar(50)	<input checked="" type="checkbox"/>
	city_name	nvarchar(50)	<input checked="" type="checkbox"/>
	bike_lanes_km	decimal(18, 10)	<input type="checkbox"/>
	protected_bike_lanes_km	decimal(18, 10)	<input type="checkbox"/>
			<input type="checkbox"/>

Figure 11 Dimension Territory Table

Dimension Store

The store table functions as a monitoring and managing the individual stores containing the relevant information about the store whilst also showcasing the costs associated with operating store (salaries and rent). It describes the scope and size of the store by rating the showcase quality. These dimensions greatly aid the company in making decisions for individual stores.

	Column Name	Data Type	Allow Nulls
▼	sk_storeid	int	<input type="checkbox"/>
	bk_store	nvarchar(50)	<input checked="" type="checkbox"/>
	area_m2	decimal(18, 10)	<input type="checkbox"/>
	rent_per_month	decimal(18, 10)	<input checked="" type="checkbox"/>
	staff_instore	int	<input type="checkbox"/>
	expert_staff	int	<input type="checkbox"/>
	average_staff_salary_month	int	<input type="checkbox"/>
	accessibility_1to5	int	<input type="checkbox"/>
	showcase_quality_level_1to5	int	<input type="checkbox"/>
	active	int	<input type="checkbox"/>

Figure 4: Dimension Store Table

Dimension Customer

The customer table describes the client who places the order with the dimensions that are needed to complete the order whilst describing the client itself (numbers of cars owned, number of children) which provides valuable information for future marketing strategies.

	Column Name	Data Type	Allow Nulls
▼	sk_customerid	int	<input type="checkbox"/>
	bk_customer	int	<input type="checkbox"/>
	customer_name	nvarchar(50)	<input checked="" type="checkbox"/>
	birth_date	date	<input type="checkbox"/>
	marital_status	nvarchar(50)	<input checked="" type="checkbox"/>
	gender	nvarchar(50)	<input checked="" type="checkbox"/>
	yearly_income	int	<input type="checkbox"/>
	number_children_at_home	int	<input type="checkbox"/>
	occupation	nvarchar(50)	<input checked="" type="checkbox"/>
	number_cars_owned	int	<input type="checkbox"/>
			<input type="checkbox"/>

Figure 5: Dimension Customer Table

Dimension Sales

The fact sales table is the most comprehensive table in the database as it provides most of the actionable evidence in terms of analysing sales and profit. This table has foreign keys that link it to other tables which paints the full picture of the company’s growth/decline.

	Column Name	Data Type	Allow Nulls
🔗	fk_date	int	<input type="checkbox"/>
🔗	fk_productid	int	<input type="checkbox"/>
🔗	fk_customerid	int	<input type="checkbox"/>
🔗	fk_storeid	int	<input type="checkbox"/>
🔗	fk_territoryid	int	<input type="checkbox"/>
	sales_order_number	nvarchar(100)	<input checked="" type="checkbox"/>
	order_quantity	int	<input type="checkbox"/>
	unit_price	decimal(15, 3)	<input type="checkbox"/>
	product_standard_cost	decimal(10, 3)	<input type="checkbox"/>
	sales_amount	decimal(15, 5)	<input type="checkbox"/>
	tax_amount_unit	decimal(10, 3)	<input type="checkbox"/>
	tax_amount	decimal(15, 3)	<input type="checkbox"/>
	freight	decimal(10, 5)	<input type="checkbox"/>
	total_cost	decimal(15, 5)	<input type="checkbox"/>

Figure 24 Dimension Sales Table

Data warehouse schema

We employed the Kimball-methodology for denormalizing data to optimize the steps of queries which results in simplified monitoring of business reports whilst also boasting significant performance improvements. This sort of data modelling is considered the industry standard and is used by many different companies. Simplifying the data results in faster query searches and makes the data much more understandable for non-technical users. A star schema makes the data expandable for the future.

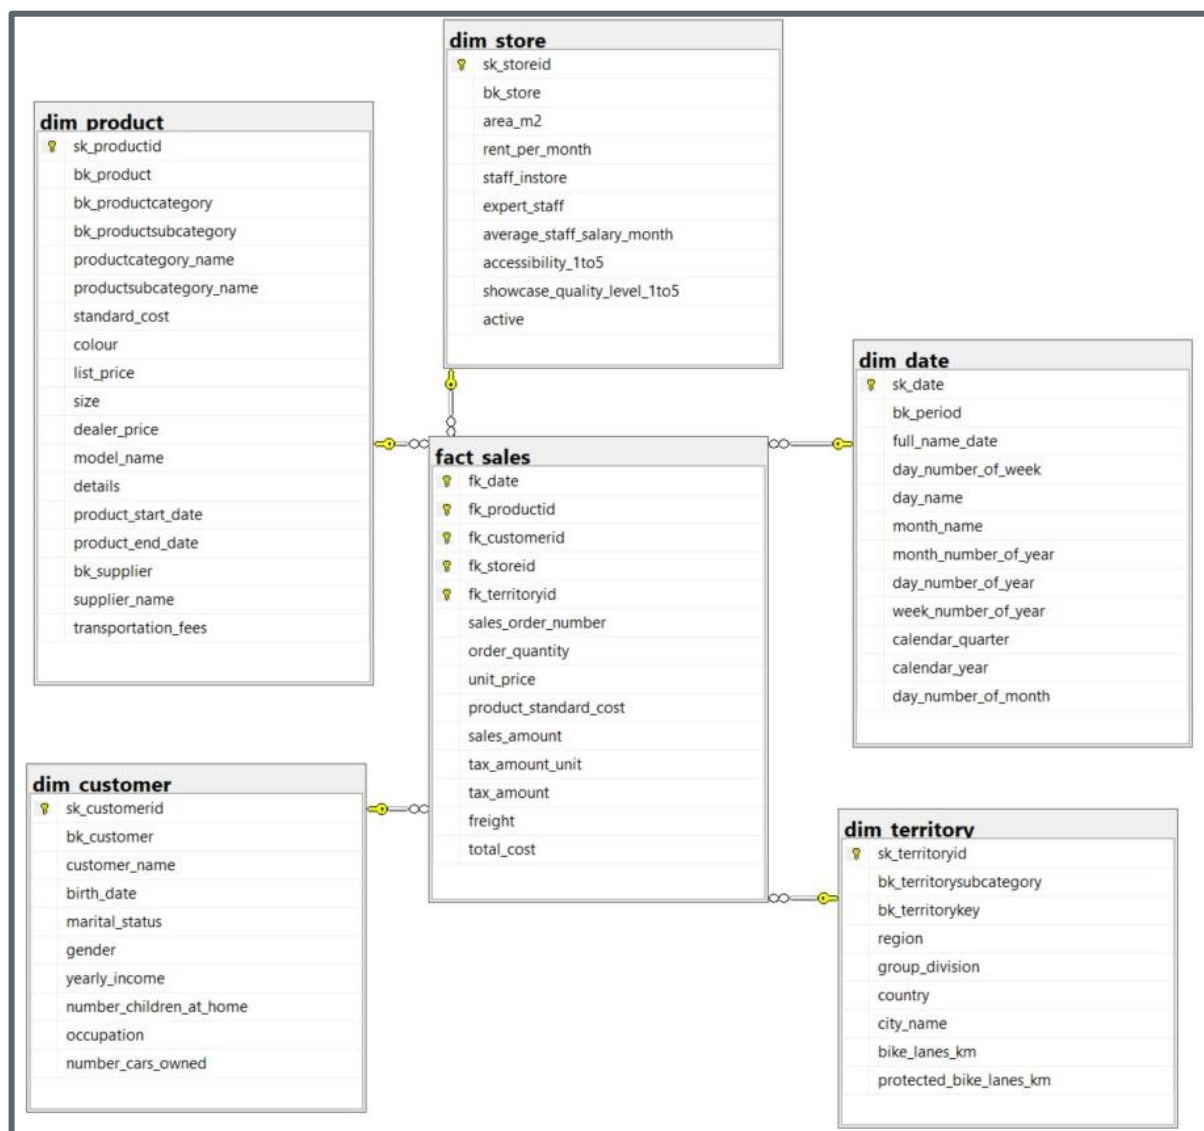


Figure 30 Data Warehouse schema

Hierarchies

Hierarchies are maps for data that go from the lowest level to the highest level and typically range from three to five levels. The creation of these maps helps on searching and analysing the data in a more optimal way. At KGB Data Warehouse we set these hierarchies:

- Date Hierarchy
- Product Hierarchy
- Customer Hierarchy
- Store Hierarchy
- Territory Hierarchy

Date Hierarchy

The Date dimension hierarchy will have six levels of depth. Having the Year at the highest level and the Full name date at the lowest level dimension. If we consider the inconsistent levels of depth between the branches, we could say that this is an unbalanced hierarchy.

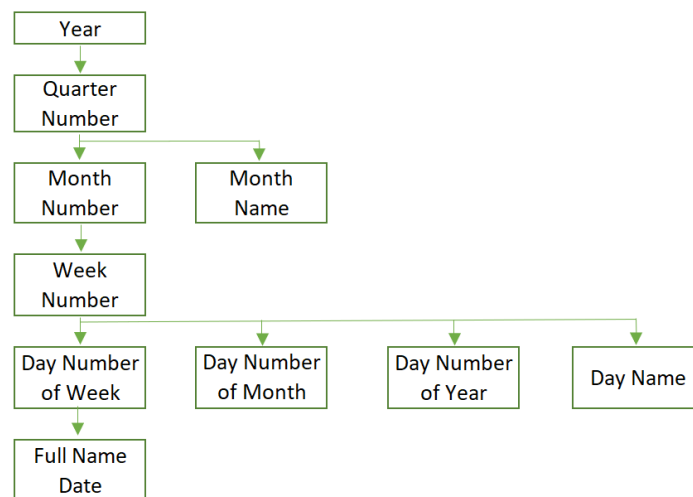


Figure 46 Date hierarchy

Product Hierarchy

For the Product dimension hierarchy will have five levels of depth. Having the Supplier at the highest level and the Model Name at the lowest level, which is divided in Price, Cost, Details, Size and Colour. We can observe in the figure below, that this is a balanced hierarchy, because the branches are descending to the same level.




Figure 16 Product hierarchy

Customer Hierarchy

The Customer dimension hierarchy will have five levels of depth. At the highest level we have Gender and at the lowest level dimension we have four different attributes that are Occupation, Birth Date, Yearly Income and Customer name. We could say for the figure below that this is a balanced hierarchy.

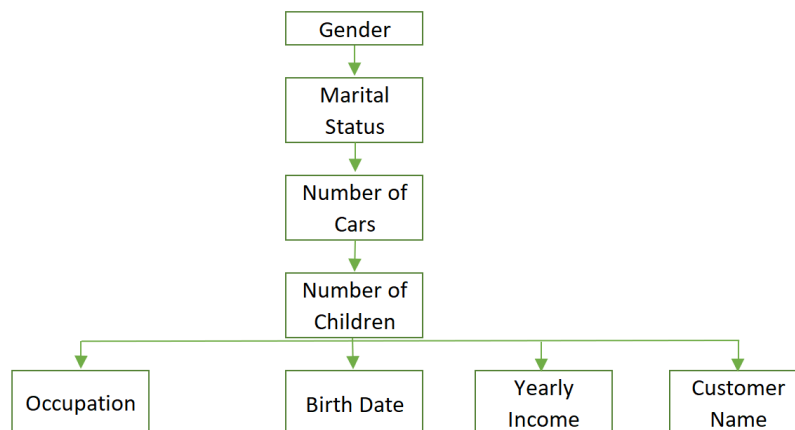


Figure 17 Customer hierarchy

Store Hierarchy

As we can see in the figure below, Store dimension hierarchy will have four levels, the highest level will be Accessibility and the lowest level dimension will be Store Code. At level three we have five different attributes that are Area, rent per Month, Staff in Store, Expert Staff and Average Salary. Because of that this is an unbalanced hierarchy.

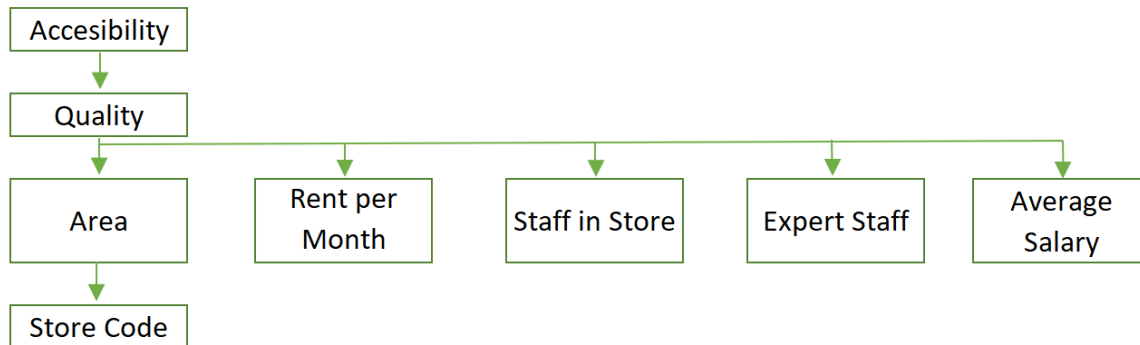


Figure 20: Store hierarchy

Territory Hierarchy

For the Territory dimension hierarchy, we will have a depth of five levels. At the highest level we have the Group and at the lowest level dimension we have City that will be divided into Bike Lanes (km) and Protected Bike Lanes (km). As we can observe in the figure below, this is a balanced hierarchy.

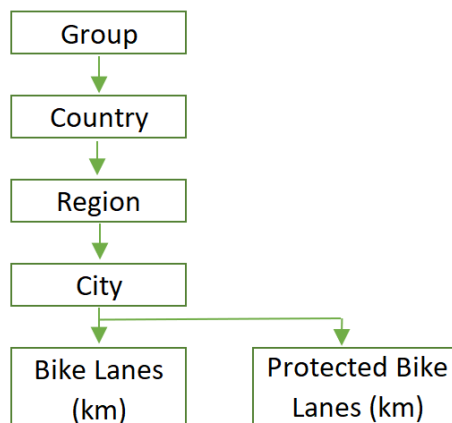


Figure 21: Territory hierarchy

Critical Changes Made After Feedback

We reevaluated the code from the first phase of the report and made changes to make the staging area and loading of the data warehouse much efficient and terminologically sound.

The changes we made and why we made them are as follows:

1. There were significant changes made to the code of the Data warehouse
 - We made sure that the dim date table had no attributes that could be null

	Column Name	Data Type	Allow Nulls
🔑	sk_date	int	<input type="checkbox"/>
	bk_period	date	<input type="checkbox"/>
	full_name_date	nvarchar(50)	<input type="checkbox"/>
	day_number_of_week	int	<input type="checkbox"/>
	day_name	nvarchar(50)	<input type="checkbox"/>
	month_name	nvarchar(50)	<input type="checkbox"/>
	month_number_of_year	int	<input type="checkbox"/>
	day_number_of_year	int	<input type="checkbox"/>
	week_number_of_year	int	<input type="checkbox"/>
	calendar_quarter	int	<input type="checkbox"/>
	calendar_year	int	<input type="checkbox"/>
	day_number_of_month	int	<input type="checkbox"/>

Figure 16 Dim_date repolished

- We looked at the product, store, territory, and customer tables and changed some attributes to null and some to not null depending on the data, and we used identity to generate surrogate keys.

	Column Name	Data Type	Allow Nulls
🔑	bk_customer	int	<input type="checkbox"/>
	customer_name	nvarchar(50)	<input checked="" type="checkbox"/>
	birth_date	date	<input checked="" type="checkbox"/>
	marital_status	nvarchar(50)	<input checked="" type="checkbox"/>
	gender	nvarchar(50)	<input checked="" type="checkbox"/>
	yearly_income	int	<input checked="" type="checkbox"/>
	number_children_at_home	int	<input checked="" type="checkbox"/>
	occupation	nvarchar(50)	<input checked="" type="checkbox"/>
	number_cars_owned	int	<input checked="" type="checkbox"/>
			<input type="checkbox"/>

Figure 17 Dim_customer repolished

	Column Name	Data Type	Allow Nulls
🔍	bk_store	nvarchar(50)	<input type="checkbox"/>
	area_m2	decimal(18, 10)	<input type="checkbox"/>
	rent_per_month	decimal(18, 10)	<input type="checkbox"/>
	staff_instore	int	<input type="checkbox"/>
	expert_staff	int	<input type="checkbox"/>
	average_staff_salary_month	int	<input type="checkbox"/>
	accessibility_1to5	int	<input type="checkbox"/>
	showcase_quality_level_1to5	int	<input type="checkbox"/>
			<input type="checkbox"/>

Figure xxi Dim_store repolished

	Column Name	Data Type	Allow Nulls
🔍	sk_territory	int	<input type="checkbox"/>
	bk_territory	nvarchar(50)	<input type="checkbox"/>
	region	nvarchar(50)	<input checked="" type="checkbox"/>
	group_division	nvarchar(50)	<input checked="" type="checkbox"/>
	country	nvarchar(50)	<input checked="" type="checkbox"/>
	city_name	nvarchar(50)	<input checked="" type="checkbox"/>
	bike_lanes_km	decimal(18, 10)	<input checked="" type="checkbox"/>
	protected_bike_lanes_km	decimal(18, 10)	<input checked="" type="checkbox"/>

Figure xxi Dim_territory repolished

	Column Name	Data Type	Allow Nulls
🔍	bk_product	int	<input type="checkbox"/>
	productcategory	nvarchar(50)	<input type="checkbox"/>
	productsubcategory	nvarchar(50)	<input type="checkbox"/>
	productcategory_name	nvarchar(50)	<input type="checkbox"/>
	productsubcategory_name	nvarchar(50)	<input type="checkbox"/>
	standard_cost	decimal(18, 10)	<input type="checkbox"/>
	colour	nvarchar(50)	<input checked="" type="checkbox"/>
	list_price	decimal(18, 10)	<input type="checkbox"/>
	size	int	<input checked="" type="checkbox"/>
	dealer_price	decimal(18, 10)	<input type="checkbox"/>
	model_name	nvarchar(100)	<input type="checkbox"/>
	details	nvarchar(350)	<input type="checkbox"/>
	product_start_date	date	<input type="checkbox"/>
	product_end_date	date	<input type="checkbox"/>
	supplier	nvarchar(50)	<input type="checkbox"/>
	supplier_name	nvarchar(50)	<input type="checkbox"/>
	transportation_fees	int	<input type="checkbox"/>

Figure xxi Dim_product repolished

- Corrected the specified constraints and added sales_order_number to the constraints.

	Column Name	Data Type	Allow Nulls
🔑	fk_date	int	<input type="checkbox"/>
🔑	fk_product	int	<input type="checkbox"/>
🔑	fk_customer	int	<input type="checkbox"/>
🔑	fk_store	int	<input type="checkbox"/>
🔑	fk_territory	int	<input type="checkbox"/>
🔑	sales_order_number	nvarchar(100)	<input type="checkbox"/>
	order_quantity	int	<input type="checkbox"/>
	unit_price	decimal(15, 3)	<input type="checkbox"/>
	product_standard_cost	decimal(10, 3)	<input type="checkbox"/>
	sales_amount	decimal(15, 5)	<input type="checkbox"/>
	tax_amount_unit	decimal(10, 3)	<input type="checkbox"/>
	tax_amount	decimal(15, 3)	<input type="checkbox"/>
	freight	decimal(10, 5)	<input type="checkbox"/>
	total_cost	decimal(15, 5)	<input type="checkbox"/>
	profit	decimal(15, 5)	<input checked="" type="checkbox"/>

Figure xxv Replicated fact_sales

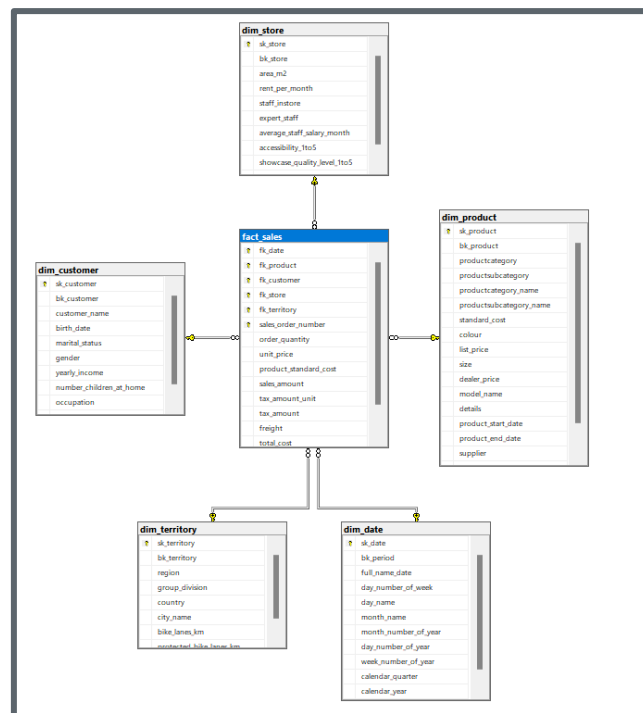


Figure xxvi Replicated Data Warehouse

2. We improved the documentation of the methodology of the established data warehouse design.

Extract, Transform and Load (ETL) Process

Our primary injection of data pertaining sales, customers, region, etc. came from SQL management studio and all information was imported as a CSV comma delimited flat file.



CSV files are commonly used in large databases because of their efficiency when it comes to read and write speeds since they are only plain text files whilst also being easy to edit. This allows for faster load times, but they also have some problems when it comes to complex data but according to our judgement, we decided to use CSV comma delimited files because our dataset felt like it could benefit from it.

ETL process was divided into 3 steps:

1. Load the data into the Staging Area.
2. Load the Staging Area to the Data Warehouse.
3. Execution of Staging Area and Data Warehouse.

First step involves the extraction of data followed by transformation. Data is then loaded into the Data Warehouse as the last step.

ETL - Staging Area

In the architecture of Data Warehouse, the staging area is where data is copied from the operating systems. It is a transformational process where data is cleansed and merged before it could be loaded into the Data Warehouse. Staging Area provides flexibility during the process of extraction and loading of data which increases the efficiency and ensures the data integrity. The design of Staging Area and Data Warehouse structure must match to achieve maximum efficiency which reduces the need for any future modifications.

To sum up the staging area is a intermediate repository where data gets itself in the best possible shape. It also:

1. Serves as a point for data correction
2. Acts as a separation from the DW because it presents the possibility of uploading it in a more controlled timeline.

We had several data types such as floats, data, currency, date, integers, strings, etc. The number of these data types had to be reduced and converted into data types like nvarchar, date, integer which are more SQL friendly. (Condensed multiple data types to 3 to improve efficiency and speed).

Methodology

The ETL staging area is a critical component of the ETL process because it allows for the separation of the extracting, transforming, and loading of the data. This separation ensures that the data is clean and accurate before it is loaded into the destination system. In addition, the ETL staging area provides a temporary storage location for the data, which can be useful when dealing with large volumes of data or when the data needs to be accessed by multiple users or processes.





In general, the ETL staging area should be designed to be scalable, secure, and efficient to ensure that it can support the needs of the organization and the ETL process.

The methodology for using a staging area in the ETL process typically involves the following steps:

1. **Extract:** The data is extracted from its source and moved to the staging area. This step may involve using extraction tools and techniques, such as SQL queries or APIs, to retrieve the data from the source system.
2. **Transform:** The data in the staging area is then transformed to fit the requirements of the destination system. This step may involve cleaning and formatting the data, as well as applying any necessary transformations or calculations.
3. **Load:** The transformed data is then loaded into the destination system. This step may involve using load tools and techniques, such as SQL INSERT statements or bulk load operations, to transfer the data from the staging area to the destination system.
4. **Verify:** After the data has been loaded into the destination system, it is typically verified to ensure that it has been loaded correctly and that it meets the requirements of the destination system. This step may involve comparing the data in the destination system to the original data in the source system or running tests to ensure that the data is accurate and complete.

By following this methodology, we can use a staging area to manage the ETL process effectively and efficiently and ensure the accuracy and integrity of the data as it is transferred from its source to the destination system.



Loading the ETL Package

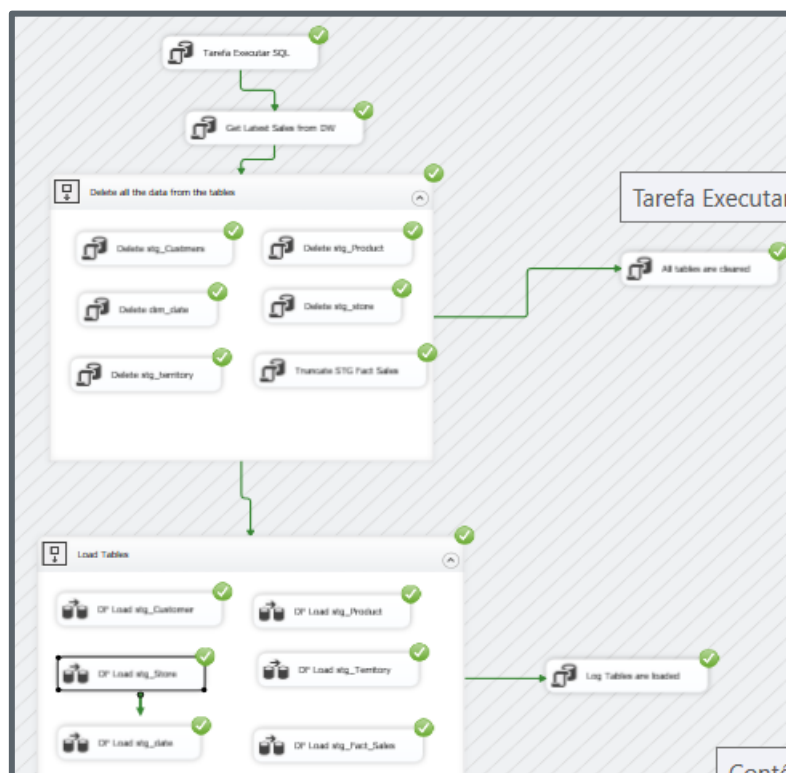


Figure xxvi Staging Area Loading - Full process

We start the ETL package by identifying the data that we wanted to extract from its source – an Excel spreadsheet. After this step, we decided to convert the data type to a CSV file.

We have logged the start runtime in the table log_stg_etl using an Execute SQL Task. We have done this by first creating two new variables, 'ETL_Name' to get the current system time, and 'ETL_Desc' to get the description of the task. These values are then stored on the Log Table to keep records of all the times the process started as shown below.

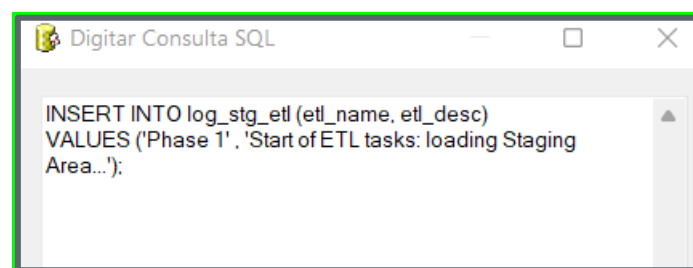


Figure xxvii One of log - stg - etl

Staging Area Table Data Clearing

The following sequence container deletes all previous data in the staging area. This is a crucial step as it helps to maintain the integrity, performance, and simplicity of the process. To clear the previous data in the staging area, we employ two methods: DELETE for the dimension tables and TRUNCATE for the fact table. As there are no relationships between the tables at this point in the project, they can all be erased simultaneously.

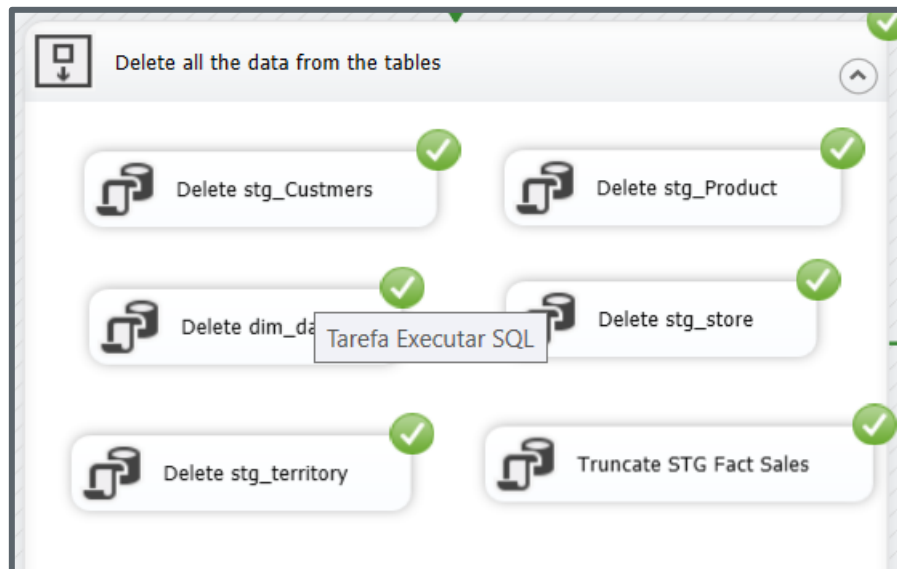


Figure 15.6: Deleting Container for Staging Area

To delete the information in the dimension tables, we use an Execute SQL Task with the code "DELETE FROM 'Dimension'". For the fact tables, we utilize the TRUNCATE function, which is a faster DDL (Data Definition Language) statement that cannot be undone, in contrast to the DELETE function which is a DML (Data Modification Language) statement that logs each row that is deleted and can be reversed. We chose to use the TRUNCATE function for the fact tables due to its ability to improve performance when deleting many rows, as was the case with the stg_fact_sales table.

We did not delete any information from the LOG_Stg_ETL table as we desire to keep a record of all the steps taken in our workflow. These records will not be transferred to the data warehouse.

Staging Area Dimensions - Data Loading

We created a new sequence container to organize all the completed loading processes for the tables in the staging area. At this stage of the project, there were no relationships between the tables, allowing us to load them simultaneously, as shown below. However, it was still important to ensure that the data was properly transformed and cleaned according to the specific requirements of each table. This required careful attention to detail and adherence to best practices. By fulfilling these requirements, we completed the loading processes and progressed to the next phase of the project.

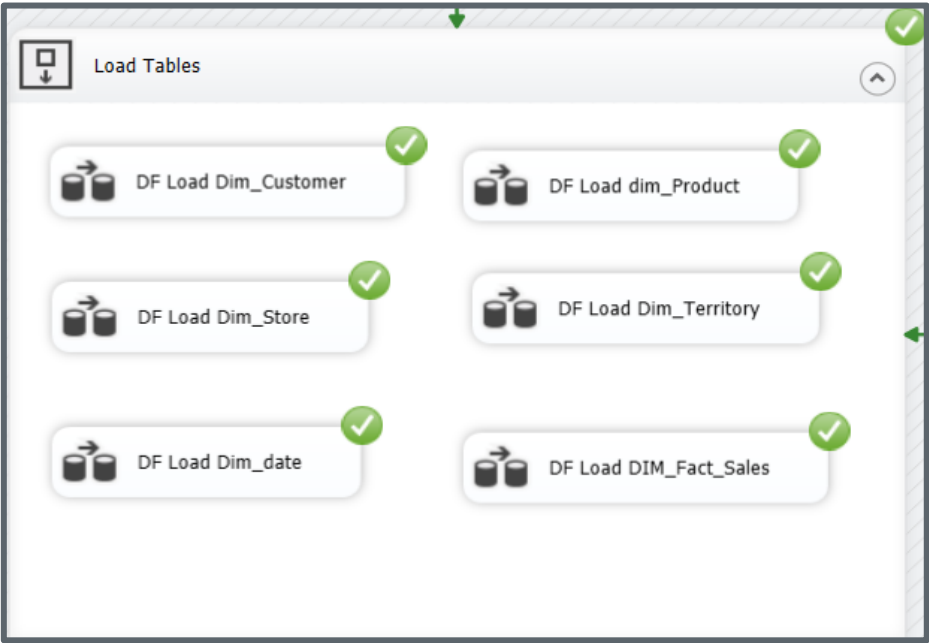


Figure xxx Loading Container for Staging Area

Loading stg_Customer

The process of importing the KGB's customer dimension was made easier by the shared structure and naming conventions between the database and the KGB's table. Mapping the input columns, as shown below, with the destination columns is simple as they all share the same name.

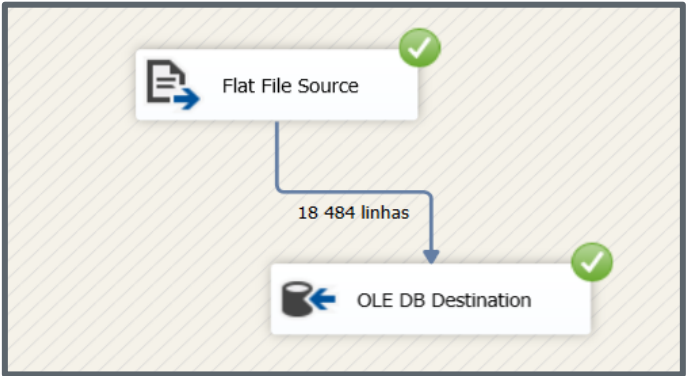


Figure xxx Loading stg_Customer



Loading stg_Product

The process of loading the KGB's product dimension into the system was a Type 0 process. One of the challenges we faced during this process was mapping the input columns, as shown below. There were some flaws in the data type conversions that needed to be addressed, which required careful attention to detail and a thorough understanding of the data to ensure that the columns were properly aligned, and any errors were corrected. Despite these challenges, we were able to successfully load the product dimension and move on to the next phase of the project.

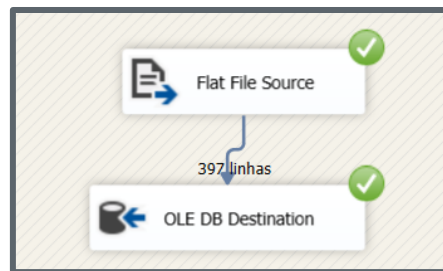


Figure xxiii Loading stg_Product

Loading stg_Store

Loading the KGB's Store dimension was also Type 0 process. Mapping the input columns, as shown below, was challenging once there were some flaws in the data types conversions.

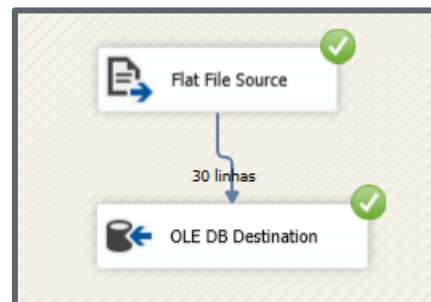


Figure xxiv Loading stg_Store

Loading stg_Territory

Loading the KGB's Territory dimension was also a Type 0 process.

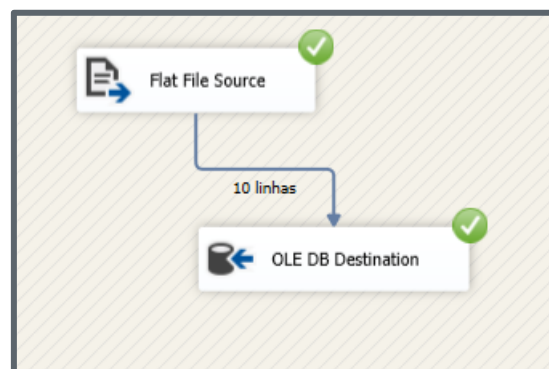


Figure xxvii Loading stg_Territory



Loading stg_Date

Loading the KGB's Date dimension was also a Type 0 process.

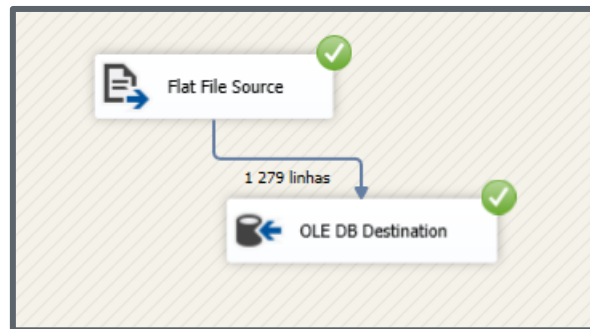


Figure 3346 Loading stg_Date

Loading Dim_Fact_Sales

Loading the KGB's Fact Sales staging dimension was a different process: we had to ensure ways of doing the desired incremental loading. We achieved this by the usage of a *Conditional Split* after applying a *Data Conversion*.

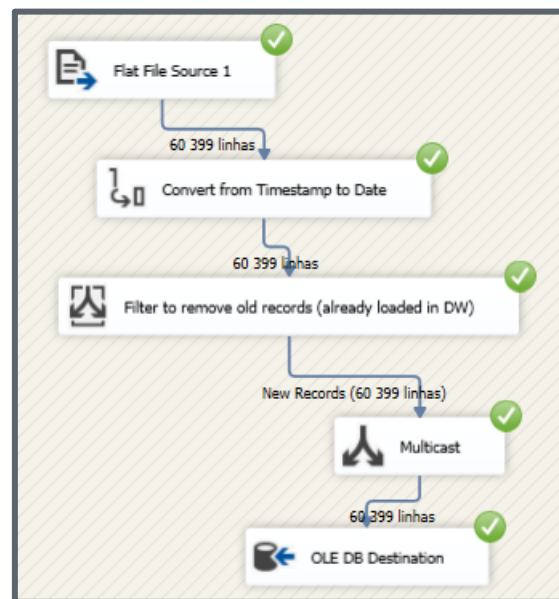


Figure 3347 stg_fact_sales loading

ETL – Loading the Data Warehouse

At the time the data is polished and reshaped within the Staging Area, the process of loading it into the Data Warehouse may go ahead. This stage of the assembly of the Business Intelligence project is where data truly becomes information. Although it had already gone through the procedures described in the previous chapter, only now the knowledge will be properly stored so it can be a resourceful helping tool to KGB: the Data Warehouse would now be able to start working. Undesirably, we were not able to properly load the *fact_table*. We will later cover this topic in the conclusion.

In this segment we describe all the operations that took place to get the data loaded to the rest of the data warehouse tables.

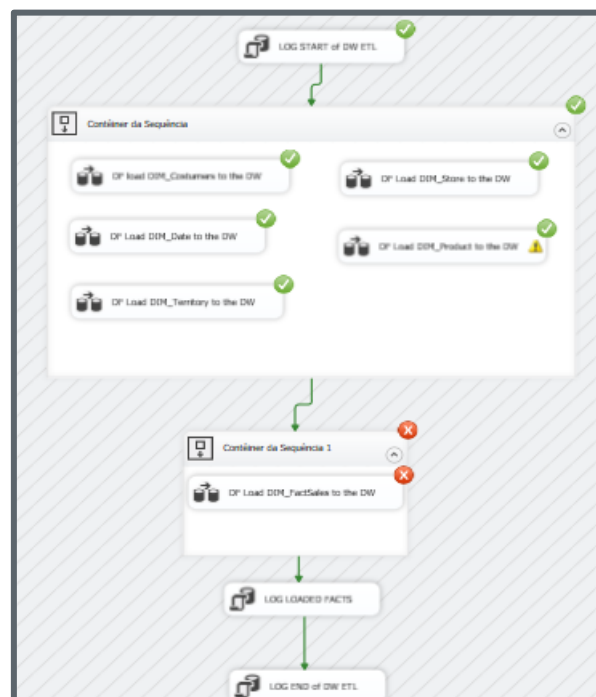



Figure 33: Data Warehouse Load

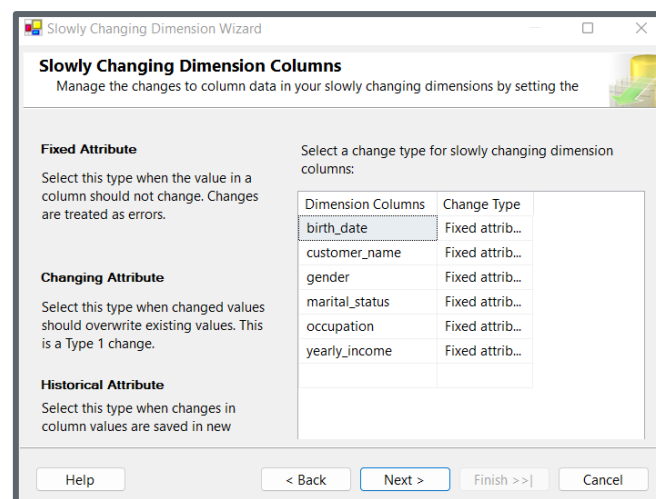
Approach on the Slowly Changing Dimensions (SCD)

To begin, we gathered information from the Staging Area and linked each Data Warehouse dimension with its corresponding dimension in the SA for all dimension tables. After this information was collected, we set up Slowly Changing Dimensions (SCD) for each dimension and specified the SCD types for each attribute. If we are performing an incremental data load, it is necessary to ensure that any existing data is updated without causing an error (type 1 SCD) while also keeping track of these changes (type 2 SCD).



Loading Dim_Customer

Regarding the Dim_Customer data flow, we've configured the birth_date, customer_name, gender, marital_status, occupation and yearly_income attributes as SCD type 0 (Fixed Attribute), because we believe that new data on these attributes should overwrite the existing data in the dimension table with new data from the source system. We have so on considered that the source system is the primary source of truth for the data and any changes made to the data in the dimension table should be overwritten by the source system. We also believe that the data in the source system is clean and does not require any further transformation or cleansing before being loaded into the dimension table.



Source: [Excel Loading Customer](#)

Loading Dim_Product

Looking at the Products Dimension, our decision was to overwrite existing values in supplier_name (making it a type 1 SCD) and to save all the values regarding dealer_price, standard_cost and transportation_fees by defining the change type as SCD 2 (Historical Attribute). We consider that it's relevant to keep historical track of changes in these dimensions as shown below.



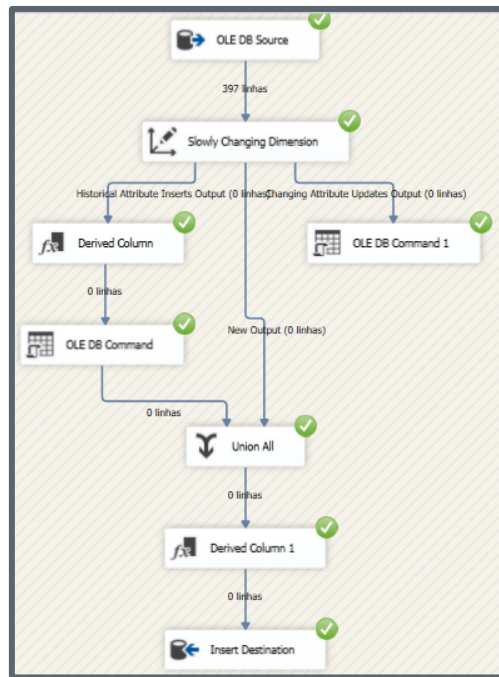


Figure 30: Loading Dim_Product

Loading Dim_Store

Regarding the Dim_Store data flow, we've configured the accessibility_1to5, active, area_m2, average_staff_salary, expert_staff, rent_per_month and staff_instore attributes as SCD type 0 (Fixed Attribute), because we believe that the data in the source system is clean and does not require any further transformation or cleansing before being loaded into the dimension table.

Loading Dim_Date

Regarding the Dim_Date data flow, we've configured all attributes as SCD type 0 (Fixed Attribute).

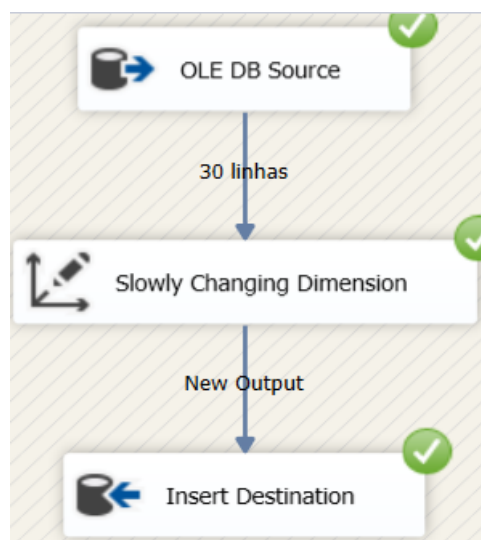


Figure 31: Loading Dim_Store

Regarding the Dim_Territory data flow, we've configured all attributes as SCD type 0 (Fixed Attribute).

Dimension Columns	Change Type
city_name	Fixed attrib...
country	Fixed attrib...
group_division	Fixed attrib...
region	Fixed attrib...
territorysubcategory	Fixed attrib...

Figure xl Loading Dim Territory

Loading Fact Sales

In the process of loading Fact Sales, we could not achieve the desired final outcome: the full load of the Fact Sales dimension.

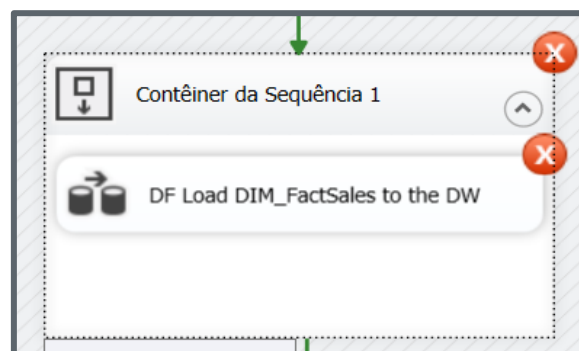


Figure xli Loading Fact Sales

Conclusion

KGB had some challenges with the way their data was stored and had severe shortcomings in optimization. After analysing the data and assessing the needs in the context of the business, we designed a solution for a Data Warehouse.

We extensively went over the literature provided to us during our classes to question the validity of the data warehouses employed and after significant amount of troubleshooting, we reached a significant point where we were satisfied by our design implementation of our Data warehouse and ETL processes.

As there are multiple ways to solve a problem when it comes to the transformational side of data handling, we had to rely on a trial-and-error methodology to assess which solution was the most fitting for our design. The staging area proved to be very beneficial and showcased



its true importance before the data was loaded into the warehouse because of the abundance of data types and formats.

Concerning to a business point of view:

In business intelligence the dimensional model is way of organizing data that makes it easy to understand and analyze. One key aspect of the dimensional model is the use of dimensions, which are attributes or categories that can be used to slice and dice the data in different ways.

The inflection points of most queries in a business-related system are related to sales and we specifically worked on getting that streamlined. We opted for 4 dimensions that all lead to shining a light on what we consider to be the most valuable insight according to our aforementioned business questions. Customers, store, territory, and product accentuate the sales data.

For instance, we recalled the business questions related to the biggest sales by market and date; the fact table should be able to access the customers info to retrieve the birth date. The same stands for “getting to know our client profile by product, product category and product sub-category”.

We must also remember that the data’s natural granularity obeys a certain order of thought, and we must not stray from it too much, to keep the model’s interpretation to the user.

Lastly, date dimension in a business intelligence dimensional model is fixed and can be populated with a set of valued dates, allowing for an easier pre-aggregation and faster querying within the data warehouse. This also emanates the need to check the validity of dates during processing. For that reason, we decided to keep that separate from the rest of the dimensions.

Analytic value proposition for the business

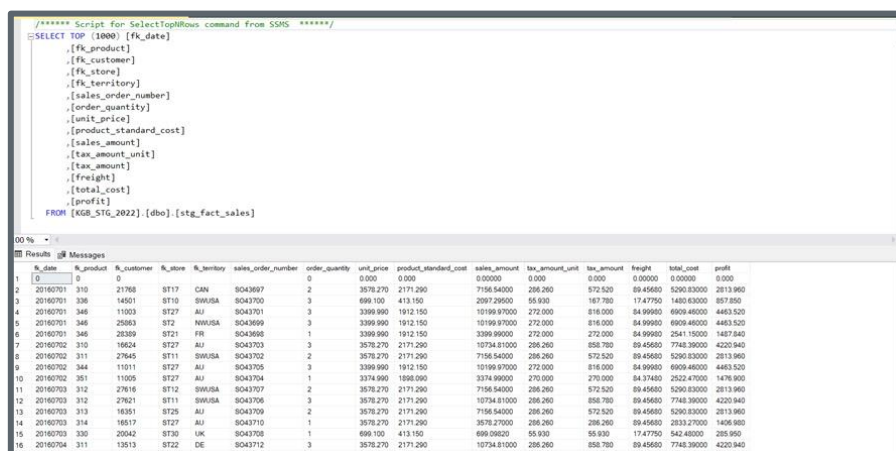
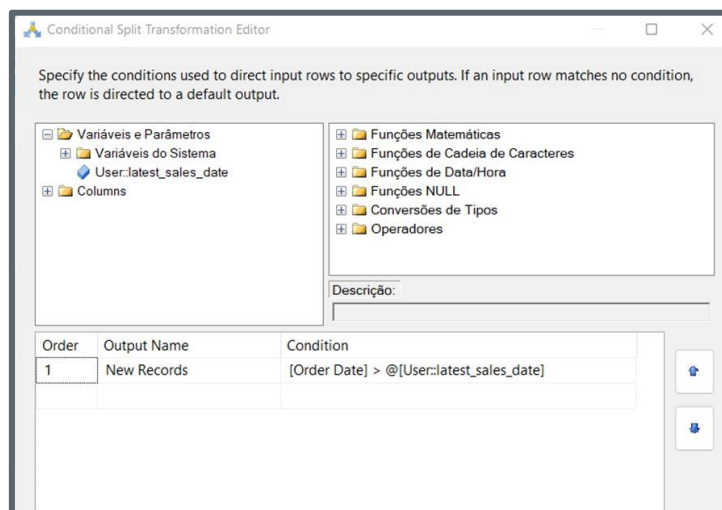
1. Improve data access – By combining all our different source files (flat files), we managed to centralize and integrate numerous aspects of KGB’s business, aiding decision makers in streamlining the process of knowledge acquirement.
2. Increase Data Quality – Through the implementation of the staging area, KGB analysts will be much more certain of dealing with meaningful data whilst also reducing the risk of errors in their analysis.
3. Better Performance – The star schema optimizes the coverage of all business aspects. For this reason, the data warehouse is optimized for reporting and analysis which means that queries will run faster, and assistant can handle larger volumes of data. Merging different dimensions to provide a bigger overall picture whilst optimizing the speed as well.
4. Improve analysis – By bringing up all the different business perspective and wrapping them in the data warehouse, KGB will be able to expand its scope on business behaviour.



5. Historical analysis – Through continuous data storage, KGB will be able to save the history of its sales. This can be useful for trend analysis, forecasting and identifying patterns in the data.
6. Data Governance – In the advent of bid data, cybersecurity plays a vital role in safeguarding the business' digital presence. Data warehousing provides a centralized system for managing and enforcing data security, compliance, and government policies, which is particularly important when dealing with sensitive data. In respects to KGB, there is a lot of data pertaining to customers that needs protective measures which is greatly improved through the implementation of data warehouses.

Critical Assessment

We were having trouble setting up the variable called “latest_sales_date” into the system variables. We searched online as to how we could solve this problem. As it is an important step in the incremental load automation part of the ETL process. We were then able to troubleshoot the problem and fixed the error.



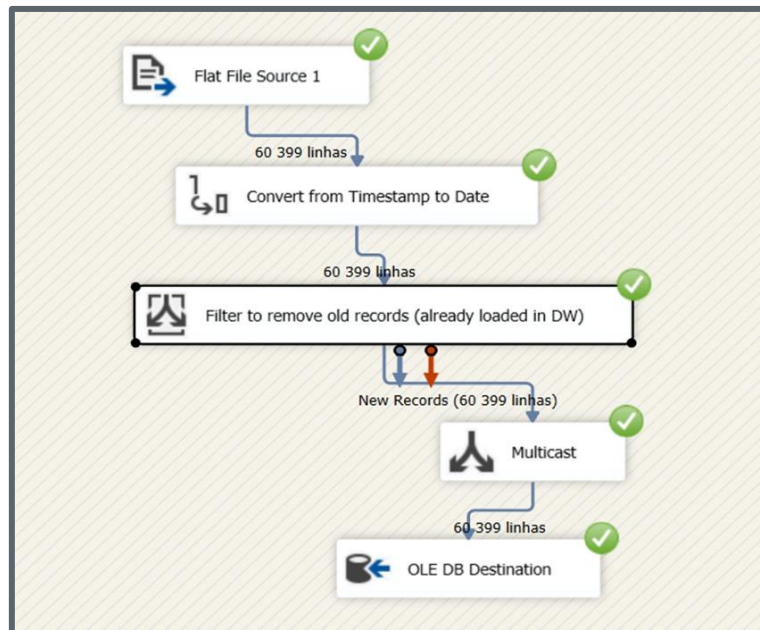


Figure 401 Error 1 Figure 1

We got this error when loading the product dimension from the fact sales table .

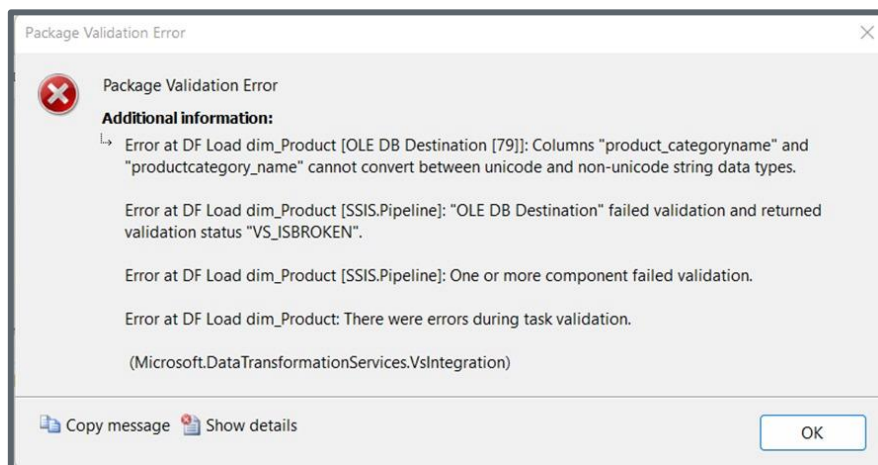


Figure 401 Error 2 Figure 2

We were then forced to change the data type by going to the connection manager tab > Flat File Connection Manager Editor > Advanced and then changed the data type to Unicode String as shown below:

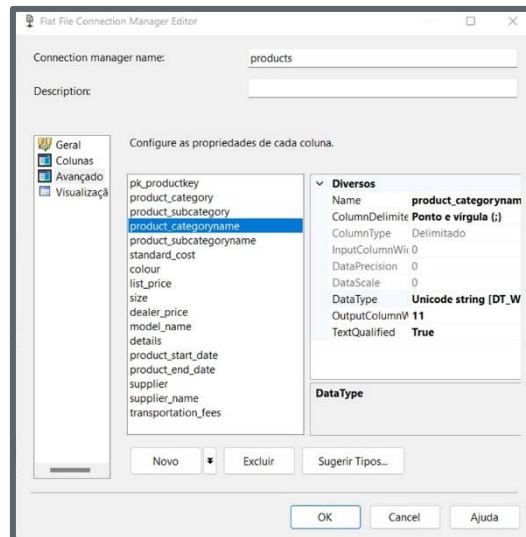


Figure 414 Error 2 Figure 4

We faced the following error when we tried loading the fact sales table from the Staging Area to the Datawarehouse.

“[Lookup Costumer [2]] Error: Row yielded no match during lookup.”

“[LookupCostumer[2]]Error:SSISErrorCode DTS_E_INDUCEDTRANSFORMFAILUREONERROR.

The "Lookup Costumer" failed because error code 0xC020901E occurred, and the error row disposition on "Lookup Costumer.Outputs[Lookup Match Output]" specifies failure on error. An error occurred on the specified object of the specified component. There may be error messages posted before this with more information about the failure.”

We tried different approaches to solve it but with no avail. We have realized that this error message typically occurs when a SQL Server Integration Services (SSIS) package is using a Lookup transformation and the data being looked up does not have a match in the reference table. With that being said, we went through an analysis of our data, and it seemed that every key condition was set properly but the error persisted. One possible reason that we thought of was one row is not “picking” the right data and because of that, we keep getting this error repeatedly. We then tried to "Redirect rows to no match output" option in the lookup transformation to capture the rows that are not getting matched, and it gave us the following result:

With this option marked, there was yet no data loaded to the Data Warehouse. We noticed that in the end of the process only 9935 lines were having a match.

Additionally, we made sure that our input data was clean and didn’t have any null or empty value in join columns. Even after trying all the remedies we could think of, the error was not fixed.

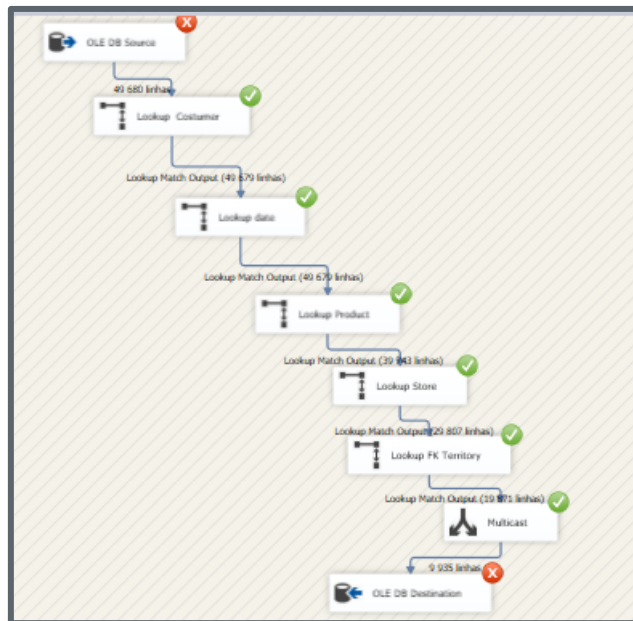


Figure 4b: Error 2